

# ndamel2-Lab4-418

November 24, 2021

## 1 Lab 4: Scikit Learn, Classification and Clustering

Deadline Tuesday 11/24/21 11:59 pm

**scikit-learn** is a popular machine learning package that contains a variety of models and tools.

All objects within scikit-learn share a uniform common basic API consisting of 3 interfaces: an *estimator* interface for building and fitting models, a *predictor* interface for making predictions, and a *transformer* interface for converting data.

The *estimator* interface defines object mechanism and a fit method for learning a model from training data. All supervised and unsupervised learning algorithms are offered as objects implementing this interface. Other machine learning tasks such as *feature extraction*, *feature selection*, and *dimensionality reduction* are provided as *estimators*.

For more information, check the scikit-learn API paper: [<https://arxiv.org/pdf/1309.0238v1.pdf>]

The general form of using models in scikit-learn:

```
clf = someModel( )  
clf.fit(x_train , y_train)
```

For Example:

```
clf = LinearSVC( )  
clf.fit(x_train , y_train)
```

The *predictor* adds a predict method that takes an array `x_test` and produces predictions for `x_test`, based on the learned parameters of the *estimator*. In supervised learning, this method typically return predicted labels or values computed by the model. Some unsupervised learning estimators may also implement the predict interface, such as **k-means**, where the predicted values are the cluster labels.

```
clf.predict(x_test)
```

*transform* method is used to modify or filter data before feeding it to a learning algorithm. It takes some new data as input and outputs a transformed version of that data. Preprocessing, feature selection, feature extraction and dimensionality reduction algorithms are all provided as *transformers* within the library.

This is usually done with **fit\_transform** method. For example:

```
PCA = RandomizedPCA (n_components = 2)  
x_train = PCA.fit_transform(x_train)
```

```
x_test = PCA.fit_transform(x_test)
```

In the example above, we first **fit** the training set to find the PC components, then they are transformed.

We can summarize the *estimator* as follows:

- In *all estimators*
  - `model.fit()` : fit training data. In supervised learning, fit will take two parameters: the data x and labels y. In unsupervised learning, fit will take a single parameter: the data x
- In *supervised estimators*
  - `model.predict()` : predict the label of new test data for the given model. Predict takes one parameter: the new test data and returns the learned label for each item in the test data
  - `model.score()` : Returns the score method for classification or regression methods.
- In *unsupervised estimators*
  - `model.transform()`: Transform new data into new basis. Transform takes one parameter: new data and returns a new representation of that data based on the model

### 1.0.1 Classification: SVM

Support Vector Machines (SVM) are among the most useful and powerful supervised learning algorithm. Here we are going to look at an example of using SVM models in scikit-learn. Then, it will be your turn to try this model.

```
[1]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
import numpy as np
```

```
[2]: from sklearn.model_selection import train_test_split
# Import make_moons from scikit learn to generate synthetic data
from sklearn.datasets import make_moons


# 2d classification dataset
Xs , ys = make_moons( n_samples = 100, noise = 0.2 , random_state = 0)

# train-test split

Xs_train , Xs_test, ys_train, ys_test = train_test_split(Xs, ys , test_size = 0.
↪15 )

#plot the data
colors = np.array(['r' , 'b'])
```

```
plt.scatter(Xs[:,0] , Xs[:,1] ,c = colors[ys] )
plt.show()
```



We will perform both linear and nonlinear SVM on this synthetic dataset:

```
[3]: def meshGrid (x , y , h):
      '''x is data for x-axis meshgrid
          y is data for y-axis meshgrid
          h is stepsize
      '''
      x_min, x_max = x.min() - 1 , x.max() + 1
      y_min, y_max = y.min() - 1 , y.max() + 1
      xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

      return xx , yy
```

```
[ ]:
```

```
[4]: #Import SVM
      from sklearn import svm

      from matplotlib.colors import ListedColormap
      from sklearn import metrics

      cmap_light = ListedColormap(['#FBBB9', '#5EFB6E', '#82CAFF'])
      cmap_bold = ListedColormap(['#CA226B', '#387C44', '#2B65EC'])
```

```

cmap_test = ListedColormap(['#8E35EF', '#FFFF00', '#659EC7'])
# cmap_light = ListedColormap(['#FBBB9', '#82CAFF'])
# cmap_bold = ListedColormap(['#CA226B', '#2B65EC'])
# cmap_test = ListedColormap(['#8E35EF', '#659EC7'])
cmap_predict = ListedColormap(['#FCDFFF', '#EOFFFF'])

# clf1 is a linear svm classifier
clf1 = svm.SVC(kernel = 'linear')

# Fit data
clf1.fit(Xs_train, ys_train)

# Predict
ys_predict = clf1.predict(Xs_test)

#Display the outcome of classification
print(metrics.classification_report(ys_test, ys_predict))
print(metrics.confusion_matrix(ys_test, ys_predict))

# Display the sum
xx , yy = meshGrid(Xs[:,0], Xs[:,1], 0.01)

Z = clf1.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.figure()
plt.contourf(xx, yy, Z, cmap=cmap_light ,levels=[-1, 0, 1] ,alpha = 0.5)

# For plotting all data use the following line
#plt.scatter(Xs[:, 0], Xs[:, 1], c=ys, cmap=cmap_bold, edgecolor='k', s=50)


# For plotting train and test and prediction separatley
plt.scatter(Xs_train[:, 0], Xs_train[:, 1], c=ys_train,
            ↪cmap=cmap_bold,edgecolor='k', s=40)
plt.scatter(Xs_test[:, 0], Xs_test[:, 1], alpha=1.0,c = ys_test,
            ↪cmap=cmap_test,linewidth=1, marker='o', s=90)
plt.scatter(Xs_test[:, 0], Xs_test[:, 1], alpha=1.0,c = ys_predict,
            ↪cmap=cmap_predict ,linewidth=1, marker='X', s=40)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())

plt.show()

```

	precision	recall	f1-score	support
0	1.00	0.89	0.94	9
1	0.86	1.00	0.92	6
accuracy			0.93	15
macro avg	0.93	0.94	0.93	15
weighted avg	0.94	0.93	0.93	15

```
[[8 1]
 [0 6]]
```



Now we apply a non-linear svm classifier

```
[5]: # clf2 is a nonlinear svm classifier

clf2 = svm.SVC(kernel = 'rbf')

# Fit data
clf2.fit(Xs_train, ys_train)

# Predict
ys_predict2 = clf2.predict(Xs_test)
```

```

#Display the outcome of classification
print(metrics.classification_report(ys_test, ys_predict2))
print(metrics.confusion_matrix(ys_test, ys_predict2))

# Display the sum
xx , yy = meshGrid(Xs[:,0], Xs[:,1], 0.01)

Z = clf2.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.figure()
plt.contourf(xx, yy, Z, cmap=cmap_light ,levels=[-1, 0, 1] ,alpha = 0.5)

# For plotting all data use the following line
#plt.scatter(Xs[:, 0], Xs[:, 1], c=ys, cmap=cmap_bold, edgecolor='k', s=50)

# For plotting train and test and prediction separatley
plt.scatter(Xs_train[:, 0], Xs_train[:, 1], c=ys_train,
    ↳cmap=cmap_bold,edgecolor='k', s=40)
plt.scatter(Xs_test[:, 0], Xs_test[:, 1], alpha=1.0,c = ys_test,
    ↳cmap=cmap_test,linewidth=1, marker='o', s=90)
plt.scatter(Xs_test[:, 0], Xs_test[:, 1], alpha=1.0,c = ys_predict2,
    ↳cmap=cmap_predict ,linewidth=1, marker='X', s=40)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())

plt.show()


```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	6
accuracy			1.00	15
macro avg	1.00	1.00	1.00	15
weighted avg	1.00	1.00	1.00	15

```

[[9 0]
 [0 6]]

```



### 1.0.2 SVM on Wine quality dataset

**Exercise 4.1 (30 pts)** Now it's your turn to work with SVM. The wine data set is loaded below. You can learn more about the dataset by using `datasett.DESCR`. Here, you need to work with the first two features to train your model.

- Select the first two features for your X
- Split the dataset in two sets of training and testing data. Use 80% of the data for training and 20% for testing
- Perform linear and non-linear SVM on the dataset
- Display the classification report and accuracy for both models

```
[6]: from sklearn.datasets import load_wine

Xwine_full , ywine = load_wine(return_X_y = True)

#Your code here

print(Xwine_full.shape, ywine.shape)
```

(178, 13) (178,)

Select the first two features for your X

```
[7]: Xwine = Xwine_full[:, :2] #using first two features
Xwine.shape
```


[7]: (178, 2)

Split the dataset in two sets of training and testing data. Use 80% of the data for training and 20% for testing

```
[8]: X_train , X_test, y_train, y_test = train_test_split(Xwine, ywine , test_size = 0.20, random_state = 42)
```

- plotting the data

```
[9]: # plot the data
colors = np.array(['r' , 'b' , 'g'])
plt.scatter(Xwine_full[:,0] , Xwine_full[:,1] , c = colors[ywine])
plt.show()
```



### 1.0.3 Perform linear and non-linear SVM on the dataset

Performing Linear SVM on dataset and displaying the classification report and accuracy

```
[10]: # clf3 is a linear svm classifier
clf3 = svm.SVC(kernel = 'linear')

# Fit data
clf3.fit(X_train, y_train)

# Predict
```



```

y_predict = clf3.predict(X_test)

#Display the outcome of classification
print(metrics.classification_report(y_test, y_predict))
print(metrics.confusion_matrix(y_test, y_predict))

```

	precision	recall	f1-score	support
0	0.79	0.79	0.79	14
1	0.92	0.86	0.89	14
2	0.56	0.62	0.59	8
accuracy			0.78	36
macro avg	0.75	0.76	0.75	36
weighted avg	0.79	0.78	0.78	36

```

[[11  0  3]
 [ 1 12  1]
 [ 2  1  5]]

```

**Performing Non-linear SVM on dataset and displaying the classification report and accuracy**

```

[11]: # clf4 is a nonlinear svm classifier

clf4 = svm.SVC(kernel = 'rbf')

# Fit data
clf4.fit(X_train, y_train)

# Predict
y_predict2 = clf4.predict(X_test)

#Display the outcome of classification
print(metrics.classification_report(y_test, y_predict2))
print(metrics.confusion_matrix(y_test, y_predict2))

```

	precision	recall	f1-score	support
0	0.85	0.79	0.81	14
1	0.93	0.93	0.93	14
2	0.56	0.62	0.59	8
accuracy			0.81	36
macro avg	0.78	0.78	0.78	36
weighted avg	0.81	0.81	0.81	36

```

[[11  0  3]
 [ 0 13  1]]

```

```
[ 2  1  5]]
```

**Exercise 4.2 (10 pts)** Scaling features is another step that can affect the performance of your classifier. For the wine data, scale the features using `StandardScaler` and perform linear SVM. Display the classification report and accuracy. Did scaling data affect the classifier performance?

```
[12]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

#Your code here
X_train , X_test, y_train, y_test = train_test_split(Xwine, ywine , test_size = 0.20, random_state = 42)
```

Scaling the training data

```
[13]: X_train_scaled=scaler.fit_transform(X_train)
      # X_train_scaled
```

Scaling the test data

```
[14]: X_test_scaled=scaler.transform(X_test)
      # X_test_scaled
```

Performing linear SVM and displaying the classification report and accuracy on the scaled data

```
[15]: # clf5 is a linear svm classifier
      clf5 = svm.SVC(kernel = 'linear')

      # Fit data
      clf5.fit(X_train_scaled, y_train)

      # Predict
      y_predict = clf5.predict(X_test_scaled)

      #Display the outcome of classification
      print(metrics.classification_report(y_test, y_predict))
      print(metrics.confusion_matrix(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.85	0.79	0.81	14
1	0.92	0.86	0.89	14
2	0.60	0.75	0.67	8
accuracy			0.81	36
macro avg	0.79	0.80	0.79	36
weighted avg	0.82	0.81	0.81	36

```
[[11  0  3]
 [ 1 12  1]
 [ 1  1  6]]
```

## CONCLUSION

- Scaling the data **improved** the classifier performance

**Exercise 4.3 (10 pts)** scikit-learn has many other classifiers. Pick another classifier of your choice ( KNN, DecisionTree, NaiveBayes, ...) and apply it to the wine dataset. Display the classification report and accuracy.

### Using Decision Tree Classifier

```
[16]: #Your code goes here
from sklearn.tree import DecisionTreeClassifier

X_train , X_test, y_train, y_test = train_test_split(Xwine, ywine , test_size = 0.20, random_state=42)

# clf6 is a Decision Tree Classifier
clf6 = DecisionTreeClassifier()

# Fit data
clf6.fit(X_train, y_train)

# Predict
y_predict = clf6.predict(X_test)

#Display the outcome of classification
print(metrics.classification_report(y_test, y_predict))
print(metrics.confusion_matrix(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.91	0.71	0.80	14
1	0.86	0.86	0.86	14
2	0.64	0.88	0.74	8
accuracy			0.81	36
macro avg	0.80	0.82	0.80	36
weighted avg	0.83	0.81	0.81	36

```
[[10  2  2]
 [ 0 12  2]
 [ 1  0  7]]
```

### Using Random Forest Classifier

```
[17]: from sklearn.ensemble import RandomForestClassifier

# clf7 is a Random Forest Classifier
clf7 = RandomForestClassifier()

# Fit data
clf7.fit(X_train, y_train)

# Predict
y_predict = clf7.predict(X_test)

#Display the outcome of classification
print(metrics.classification_report(y_test, y_predict))
print(metrics.confusion_matrix(y_test, y_predict))
```

	precision	recall	f1-score	support
0	1.00	0.71	0.83	14
1	0.87	0.93	0.90	14
2	0.73	1.00	0.84	8
accuracy			0.86	36
macro avg	0.86	0.88	0.86	36
weighted avg	0.89	0.86	0.86	36

```
[[10  2  2]
 [ 0 13  1]
 [ 0  0  8]]
```

#### 1.0.4 Clustering

You have already seen an example of clustering using scikit-learn in lecture. In this section, you will apply KMeans algorithm to the wine dataset.

##### Exercise 4.4 ( 30 pts)

- First choose the first two features and apply kmeans clustering.
- Display cluster evaluation metrics `homogeneity_score` and `completeness_score` (both belong to `sklearn.metrics`)
- Plot the clusters and centroids. You have the “ground truth” or labels of your data points, your plot should create a meshgrid to display the decision boundary of your model, and add the datapoints and their true labels. ( This is to observe how good your model performs on the data)

Note: For displaying decision boundaries and data points follow these steps:

1. Use `meshGrid` function to get the mesh for your attributes
2. Obtain labels for each point in mesh and reshape it. ( `Z = kmeans.predict(...)`)
3. Put the results into a color plot


- Plot the colormesh -> plt.pcolormesh
- Plot your data points -> plt.scatter
- Plot the centroids -> plt.scatter
- Set titles, x and y ranges
- plt.show()

```
[18]: from sklearn.cluster import KMeans

Xwine_full , ywine = load_wine(return_X_y = True)

Xc = Xwine_full[:, :2] #using first 2 features

colormap = np.array(['r' , 'b' , 'g'])
plt.scatter(Xc[:,0],Xc[:,1] , c = colormap[ywine])
plt.title("the data");
plt.show()
```



First choose the first two features and apply kmeans clustering.

```
[19]: # Your code here
kmeans = KMeans(n_clusters = 3, init = 'random' , random_state = 200 , verbose=
↳False).fit(Xc)
y_pred = kmeans.predict(Xc)
y_pred
```

```
[19]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2,
          1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 2,
          1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0,
          1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 2, 0, 0, 0, 2, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          2, 0, 0, 0, 0, 0, 0, 0, 2, 2, 0, 0, 2, 2, 2, 0, 0, 0, 0, 2, 0, 2,
          0, 2, 0, 0, 2, 2, 2, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 2, 0, 1, 2,
          0, 2, 2, 2, 1, 1, 2, 2, 2, 2, 1, 2, 2, 2, 1, 2, 2, 0, 1, 2, 2, 2,
          1, 2])
```

```
[20]: plt.scatter(Xc[:,0], Xc[:,1], c = colormap[y_pred])
plt.title("the clusters");
```



Display cluster evaluation metrics `homogeneity_score` and `completeness_score` (both belong to `sklearn.metrics`)

```
[21]: from sklearn.metrics import completeness_score, homogeneity_score

# Evaluating the performance
print('Completeness Score :', completeness_score(ywine, y_pred))
print('Homogeneity Score :', homogeneity_score(ywine, y_pred))
```

```
Completeness Score : 0.4080524820388843
Homogeneity Score : 0.41035077970969713
```

Plot the clusters and centroids. You have the “ground truth” or labels of your data

points, your plot should create a meshgrid to display the decision boundary of your model, and add the datapoints and their true labels. ( This is to observe how good your model performs on the data)


```
[22]: display(kmeans.labels_)
```

```
array([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2,
        1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 2,
        1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0,
        1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 2, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        2, 0, 0, 0, 0, 0, 0, 0, 2, 2, 0, 0, 2, 2, 2, 0, 0, 0, 0, 2, 0, 2,
        0, 2, 0, 0, 2, 2, 2, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 2, 0, 1, 2,
        0, 2, 2, 2, 1, 1, 2, 2, 2, 2, 1, 2, 2, 2, 1, 2, 2, 0, 1, 2, 2, 2,
        1, 2])
```

```
[23]: display(kmeans.cluster_centers_)
```

```
array([[12.21349206,  1.6531746 ],
        [13.71538462,  1.79969231],
        [13.0632    ,  3.8948    ]])
```

```
[24]: plt.scatter(Xc[:,0], Xc[:,1], c = colormap[kmeans.labels_])
plt.scatter(kmeans.cluster_centers_[0,0], kmeans.cluster_centers_[0,1] , c = 'k',
            ↪ 'k' , marker = 'X' , linewidths = 1)
plt.title("the clusters and centroids");
```

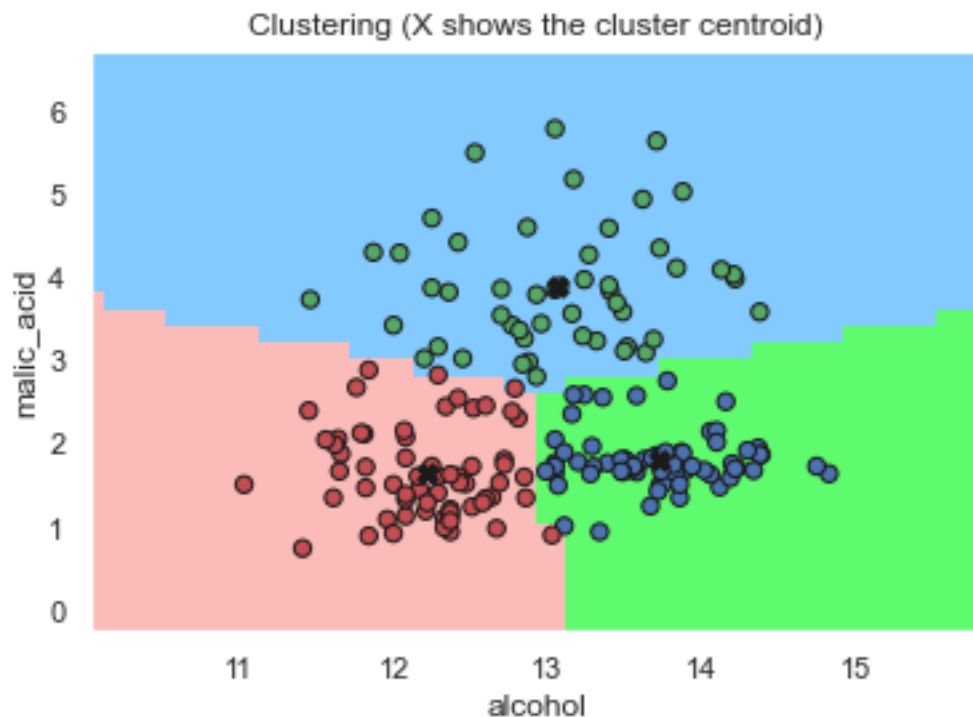


## Draw decision boundaries

```
[25]: # make the meshgrid
xx , yy = meshGrid(Xc[:,0], Xc[:,1], 0.2)

# add the classifier to the meshgrid
Z = kmeans.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# plot the outcome
plt.figure()
plt.pcolormesh(xx, yy, Z, cmap = cmap_light, shading='auto')
plt.scatter(Xc[:, 0], Xc[:, 1], c = colormap[kmeans.labels_], edgecolor = 'k', s = 40)
plt.scatter(kmeans.cluster_centers_[0,0], kmeans.cluster_centers_[0,1], c = 'k', marker = 'X', linewidths = 1, s = 55)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.xlabel("alcohol")
plt.ylabel("malic_acid")
plt.title("Clustering (X shows the cluster centroid)");
```





**Exercise 4.5 (20 pts)** In the previous model you used the first two features: ‘Alcohol’ and ‘Malic acid’. For this exercise, pick features ‘Alcohol’ and ‘OD280/OD315 of diluted wines’ (feature #1 and feature #12) as your two attributes and perform the tasks in Exercise 4.4. (cluster, report metrics, draw decision boundaries)

Which model performs better?

**Picking features ‘Alcohol’ and ‘OD280/OD315 of diluted wines’**

```
[26]: # your code here
Xc = Xwine_full[:, [0,11]] #using (feature #1 and feature #12)
# Xc
```

```
[27]: #Plotting the data
colormap = np.array(['r' , 'b' , 'g'])
plt.scatter(Xc[:,0],Xc[:,1] , c = colormap[ywine])
plt.title("the data");
plt.show()
```



**Applying kmeans clustering.**

```
[28]: # Your code here
kmeans = KMeans(n_clusters = 3, init = 'random' , random_state = 200 , verbose=
↳False).fit(Xc)
y_pred = kmeans.predict(Xc)
y_pred
```

```
[28]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
            1, 0, 2, 0, 2, 1, 1, 1, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 2, 2,
            2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
            2, 2])
```

```
[29]: plt.scatter(Xc[:,0], Xc[:,1], c = colormap[y_pred])
plt.title("the clusters");
```



Display cluster evaluation metrics `homogeneity_score` and `completeness_score`

```
[30]: from sklearn.metrics import completeness_score, homogeneity_score


# Evaluating the performance
print('Completeness Score :', completeness_score(ywine, y_pred))
print('Homogeneity Score :', homogeneity_score(ywine, y_pred))
```

Completeness Score : 0.7006853440435565

Homogeneity Score : 0.7072039236692641

Plot the clusters and centroids.

```
[31]: plt.scatter(Xc[:,0], Xc[:,1], c = colormap[kmeans.labels_], cmap = cmap_light)
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], c = 'k',
            marker = 'X', linewidths = 1)
plt.title("the clusters and centroids");
```




### Draw decision boundaries

```
[32]: # make the meshgrid
xx , yy = meshGrid(Xc[:,0], Xc[:,1], 0.2)

# add the classifier to the meshgrid
Z = kmeans.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# plot the outcome
plt.figure()
plt.pcolormesh(xx, yy, Z, cmap = cmap_light, shading='auto')
plt.scatter(Xc[:, 0], Xc[:, 1], c = colormap[kmeans.labels_], edgecolor = 'k',
            s = 40)
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], c = 'k',
            marker = 'X', linewidths = 1, s = 55)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.xlabel("alcohol")
```

```
plt.ylabel("OD280/OD315 of diluted wines")
plt.title("Clustering (X shows the cluster centroid)");
```



A clustering result satisfies homogeneity if all of its clusters contain only data points which are members of a single class.

A clustering result satisfies completeness if all the data points that are members of a given class are elements of the same cluster.

Both scores have positive values between 0.0 and 1.0, **larger values** being desirable.

## CONCLUSION

The model using features 'Alcohol' and 'OD280/OD315 of diluted wines' has a better homogeneity and completeness score than the model using the first two features and thus the model using features 'Alcohol' and 'OD280/OD315 of diluted wines' **performs better** .